

# **Model Checking for the Working Man <sup>(m/f)</sup>**

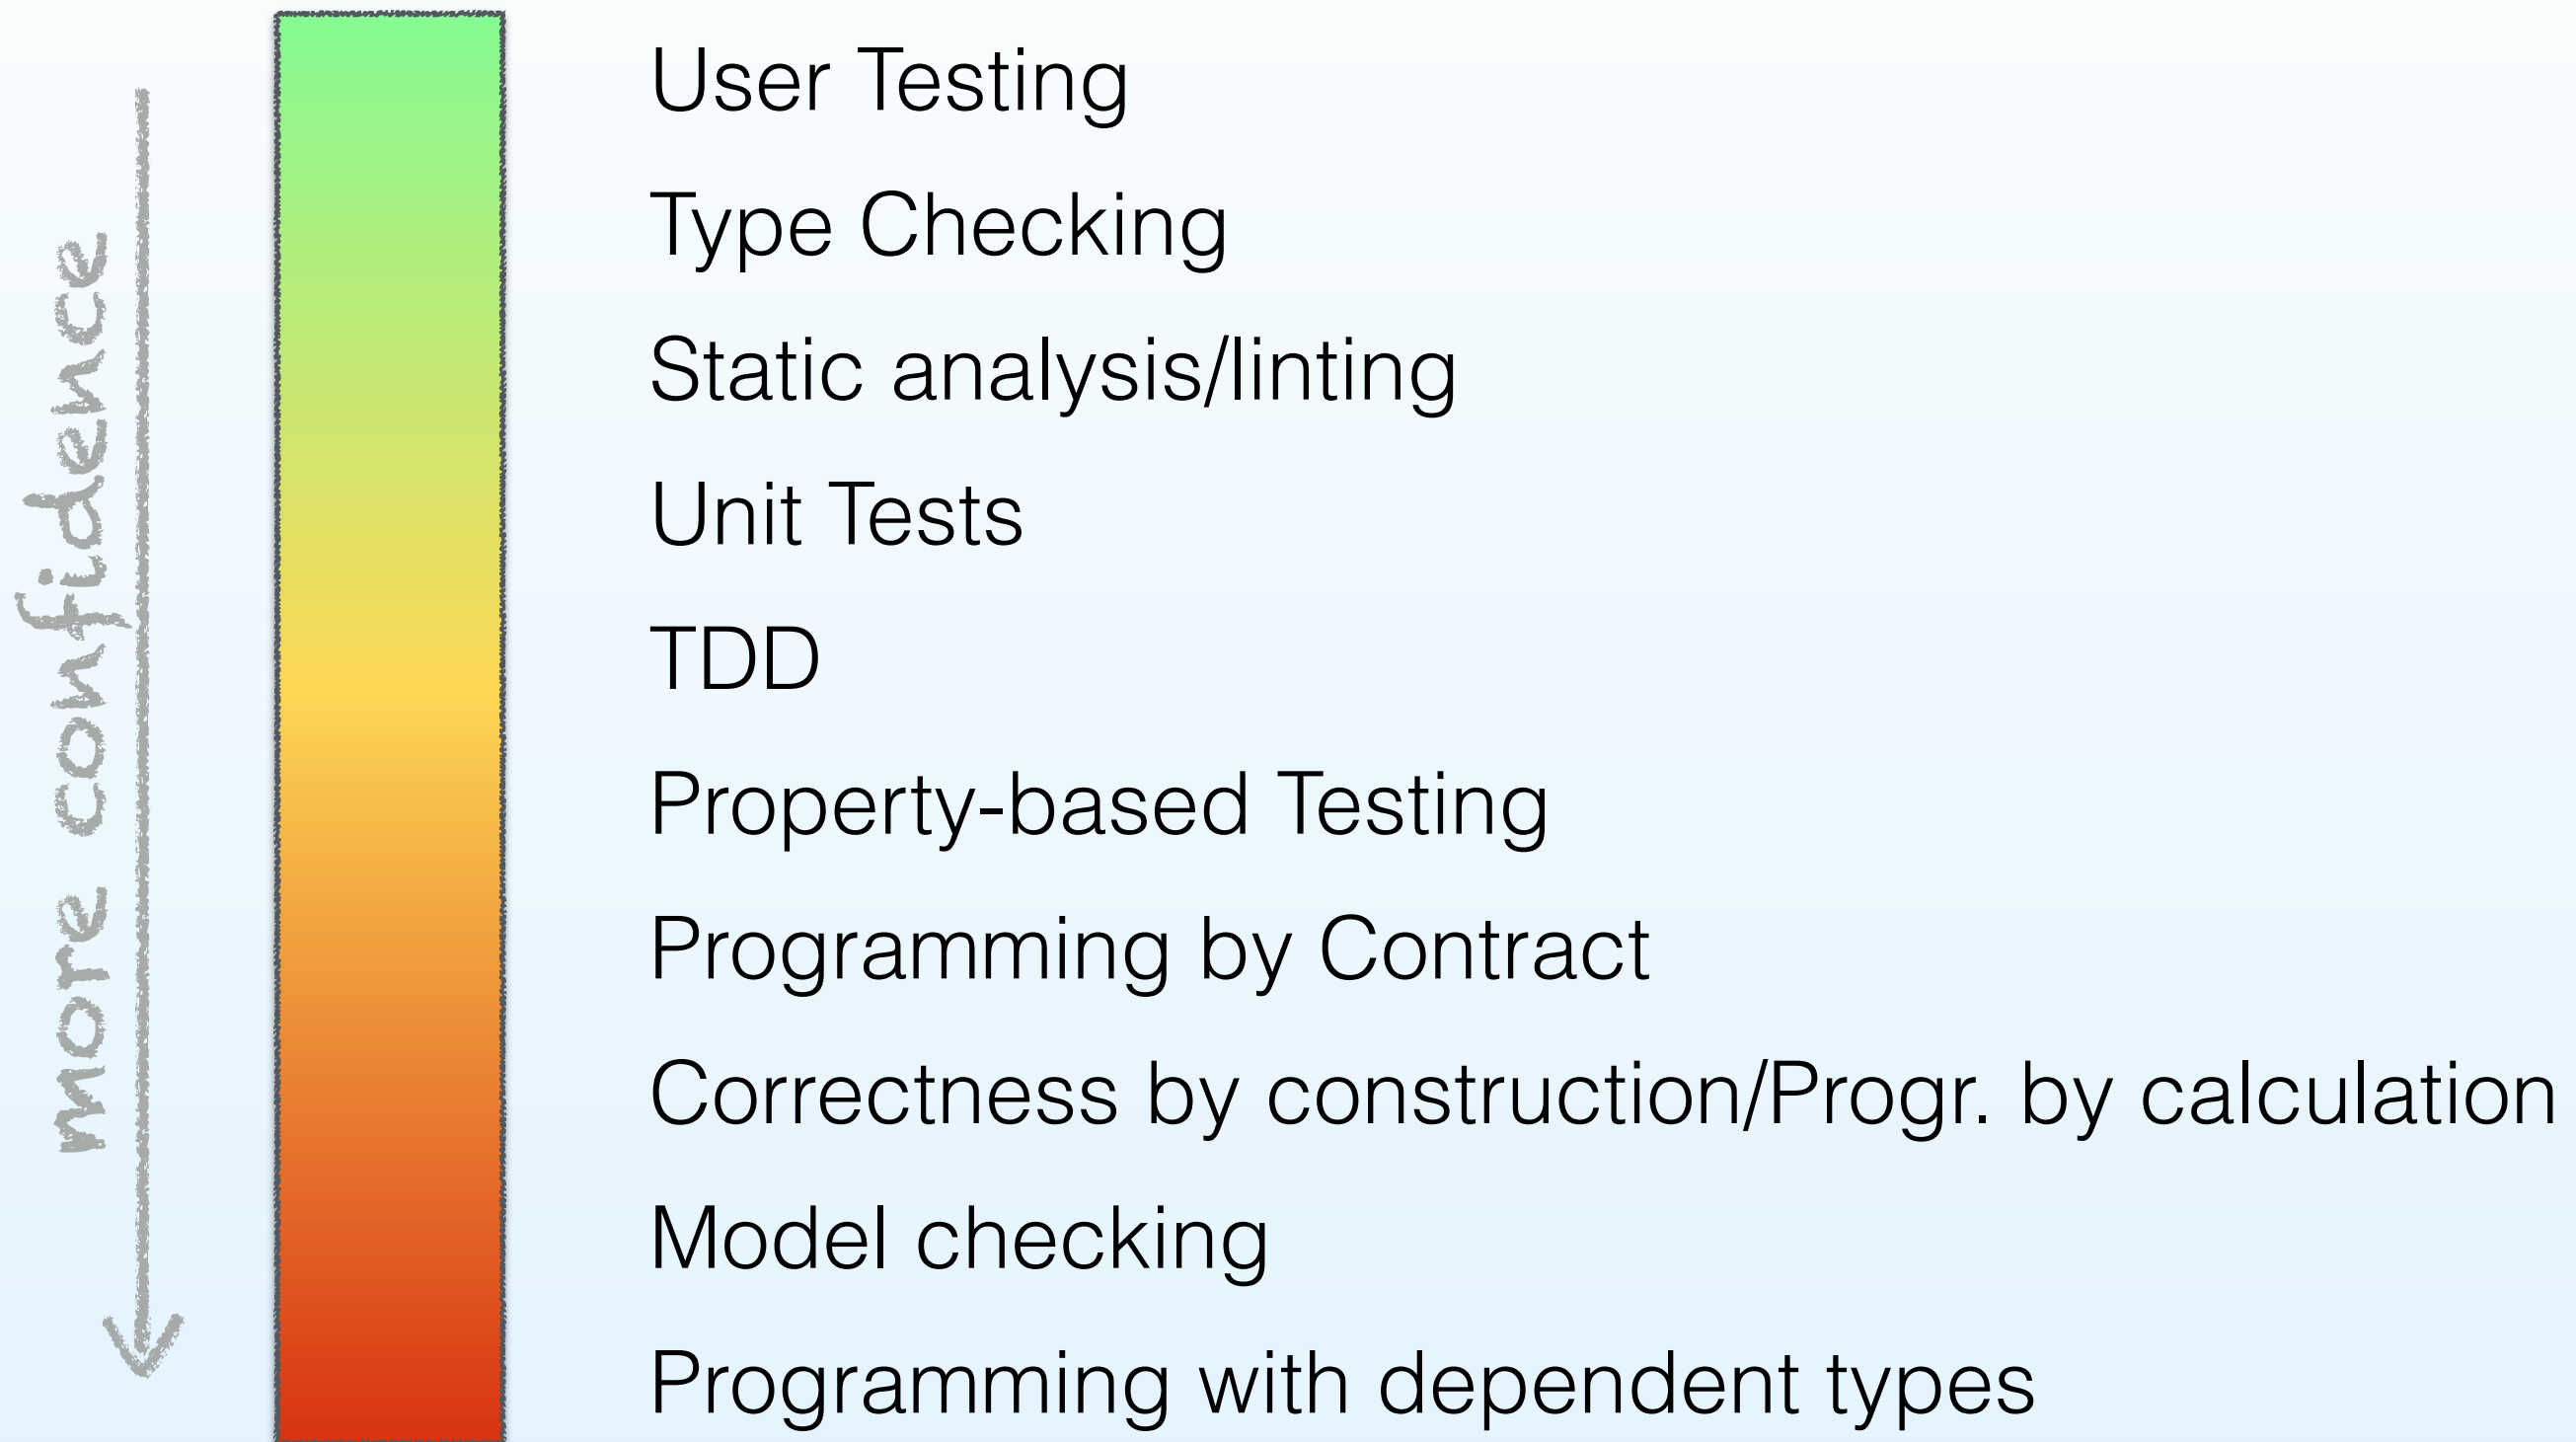
@rix0rrr



# Build the right software

- Understand the customer
- Iterative development
- Hire good programmers
- ...

# Build the software right





DECKARD CA

Stay a while  
and listen





# Model Checking

Specification

+

Model

```
x := 0;  
y := 1;  
while x < N do  
  x := x + 1;  
  y := y * 2  
end
```

$N = 10$


$\square y = 2^x$

$\diamond x = 256$



Model Checker

OK or Counterexample





**TLA+**

## Use of Formal Methods at Amazon Web Services

Chris Newcombe, Tim Rath, Fan Zhang, Bogdan Munteanu, Marc Brooker, Michael Deardeuff  
Amazon.com

29<sup>th</sup> September, 2014

Since 2011, engineers at Amazon Web Services (AWS) have been using formal specification and model checking to help solve difficult design problems in critical systems. This paper describes our motivation and experience, what has worked well in our problem domain, and what has not. When discussing personal experiences we refer to authors by their initials.

At AWS we strive to build services that are simple for customers to use. That external simplicity is built on a hidden substrate of complex distributed systems. Such complex internals are required to achieve high availability while running on cost-efficient infrastructure, and also to cope with relentless rapid business growth. As an example of this growth; in 2006 we launched S3, our Simple Storage Service. In the 6 years after launch, S3 grew to store 1 trillion objects<sup>[1]</sup>. Less than a year later it had grown to 2 trillion objects, and was regularly handling 1.1 million requests per second<sup>[2]</sup>.



*“...the TLA+ specification language [...] represents a Quixotic attempt to overcome engineers' antipathy towards mathematics.”*

# TLA+ Components

## TLA+

Temporal logic language  
*(for model, maybe spec)*

## PlusCal

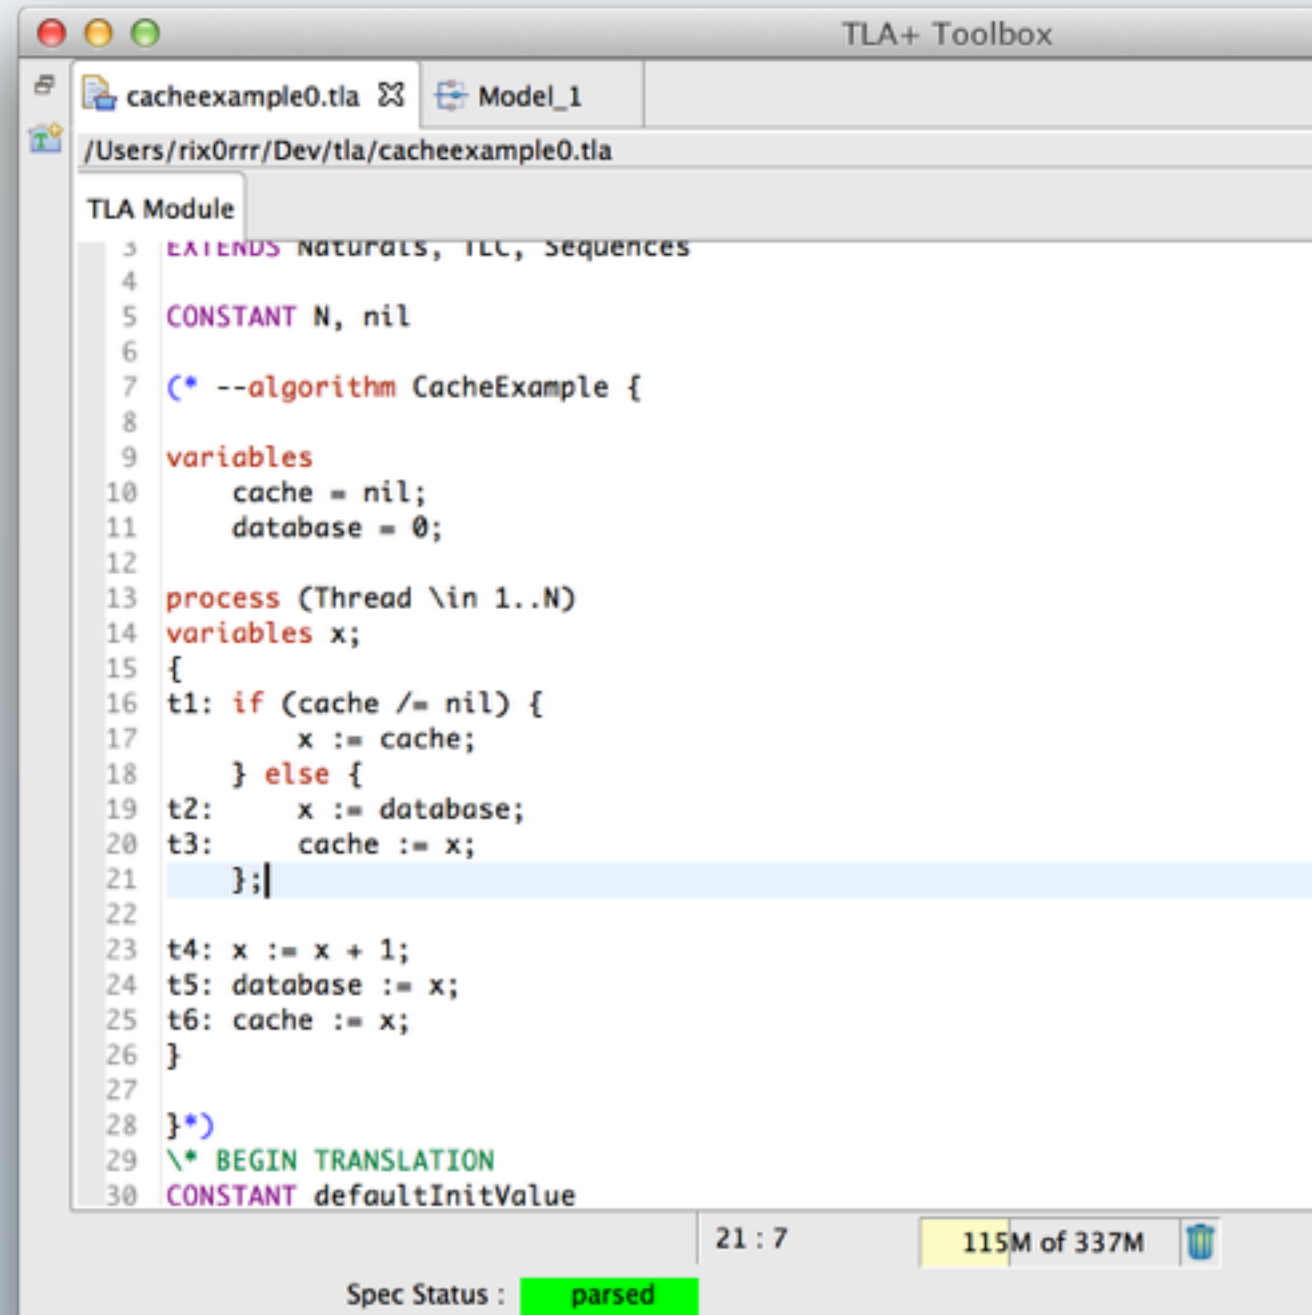
Imperative language  
*(more natural spec writing)*

## TLC

Checker tool

## Toolbox

Eclipse-based IDE



```
TLA+ Toolbox
cacheexample0.tla  Model_1
/Users/rix0rrr/Dev/tla/cacheexample0.tla

TLA Module
3  EXTENDS Naturals, TLC, Sequences
4
5  CONSTANT N, nil
6
7  (* --algorithm CacheExample {
8
9  variables
10     cache = nil;
11     database = 0;
12
13  process (Thread \in 1..N)
14  variables x;
15  {
16  t1: if (cache /= nil) {
17      x := cache;
18  } else {
19  t2:   x := database;
20  t3:   cache := x;
21  };|
22
23  t4: x := x + 1;
24  t5: database := x;
25  t6: cache := x;
26  }
27
28  }*)
29  \* BEGIN TRANSLATION
30  CONSTANT defaultInitValue

Spec Status : parsed
21 : 7
115M of 337M
```



# PlusCal

## P-syntax

```
t1:  if (x = 1) then
      begin
          x := x + 1;
          z := y;
      end;
t2:  y := x;
```

Labels indicate atomic sections


## C-syntax

```
t1:  if (x /= 2) {
          x := x + 1;
          z := y;
      };
t2:  y := x;
```

**All interleavings  
checked!**



# The Program



# A Python implementation (?)

```
cache = {}
```

```
# Database-backed incrementor  
# with cache
```

```
def Inc(id):  
    x = cache.get(id, None)  
    if x is None:  
        x = DB_Read(id)  
        cache[id] = x
```

```
    x = x + 1  
    DB_Write(id, x)  
    cache[id] = x
```

# Translating to PlusCal

```
cache = {}

# Database-backed incrementor
# with cache
def Inc(id):
    x = cache.get(id, None)
    if x is None:
        x = DB_Read(id)
        cache[id] = x

    x = x + 1
    DB_Write(id, x)
    cache[id] = x
```

Things to consider:

- How to model data? (simplify!)
- What are the atomic steps?
- What concurrency are we modelling?

# Try it!

```
cache = {}

# Database-backed incrementor
# with cache
def Inc(id):
    x = cache.get(id, None)
    if x is None:
        x = DB_Read(id)
        cache[id] = x

    x = x + 1
    DB_Write(id, x)
    cache[id] = x
```



variables

cache = nil;  
database = 0;

process (Thread \in 1..N)

variables x;

{

t1: if (cache /= nil) {  
    x := cache;  
} else {

t2: x := database;  
t3: cache := x;

};

t4: x := x + 1;

t5: database := x;

t6: cache := x;

}



# TLA+ Translation & the “pc” variable

PlusCal

```
t4: x := x + 1;
```

TLA+

```
Init == (* Global variables *)  
        /\ cache = nil  
        /\ database = 0  
        (* Process Thread *)  
        /\ x = [self \in 1..N |-> defaultInitValue]  
        /\ pc = [self \in ProcSet |-> "t1"]
```

```
t4(self) == /\ pc[self] = "t4"  
            /\ x' = [x EXCEPT ![self] = x[self] + 1]  
            /\ pc' = [pc EXCEPT ![self] = "t5"]  
            /\ UNCHANGED << cache, database >>
```

# The Condition?



"All increments should  
have been applied"

What if I tell you  
you can see if a process  
has finished by checking  
`pc[i] = "Done"` ?

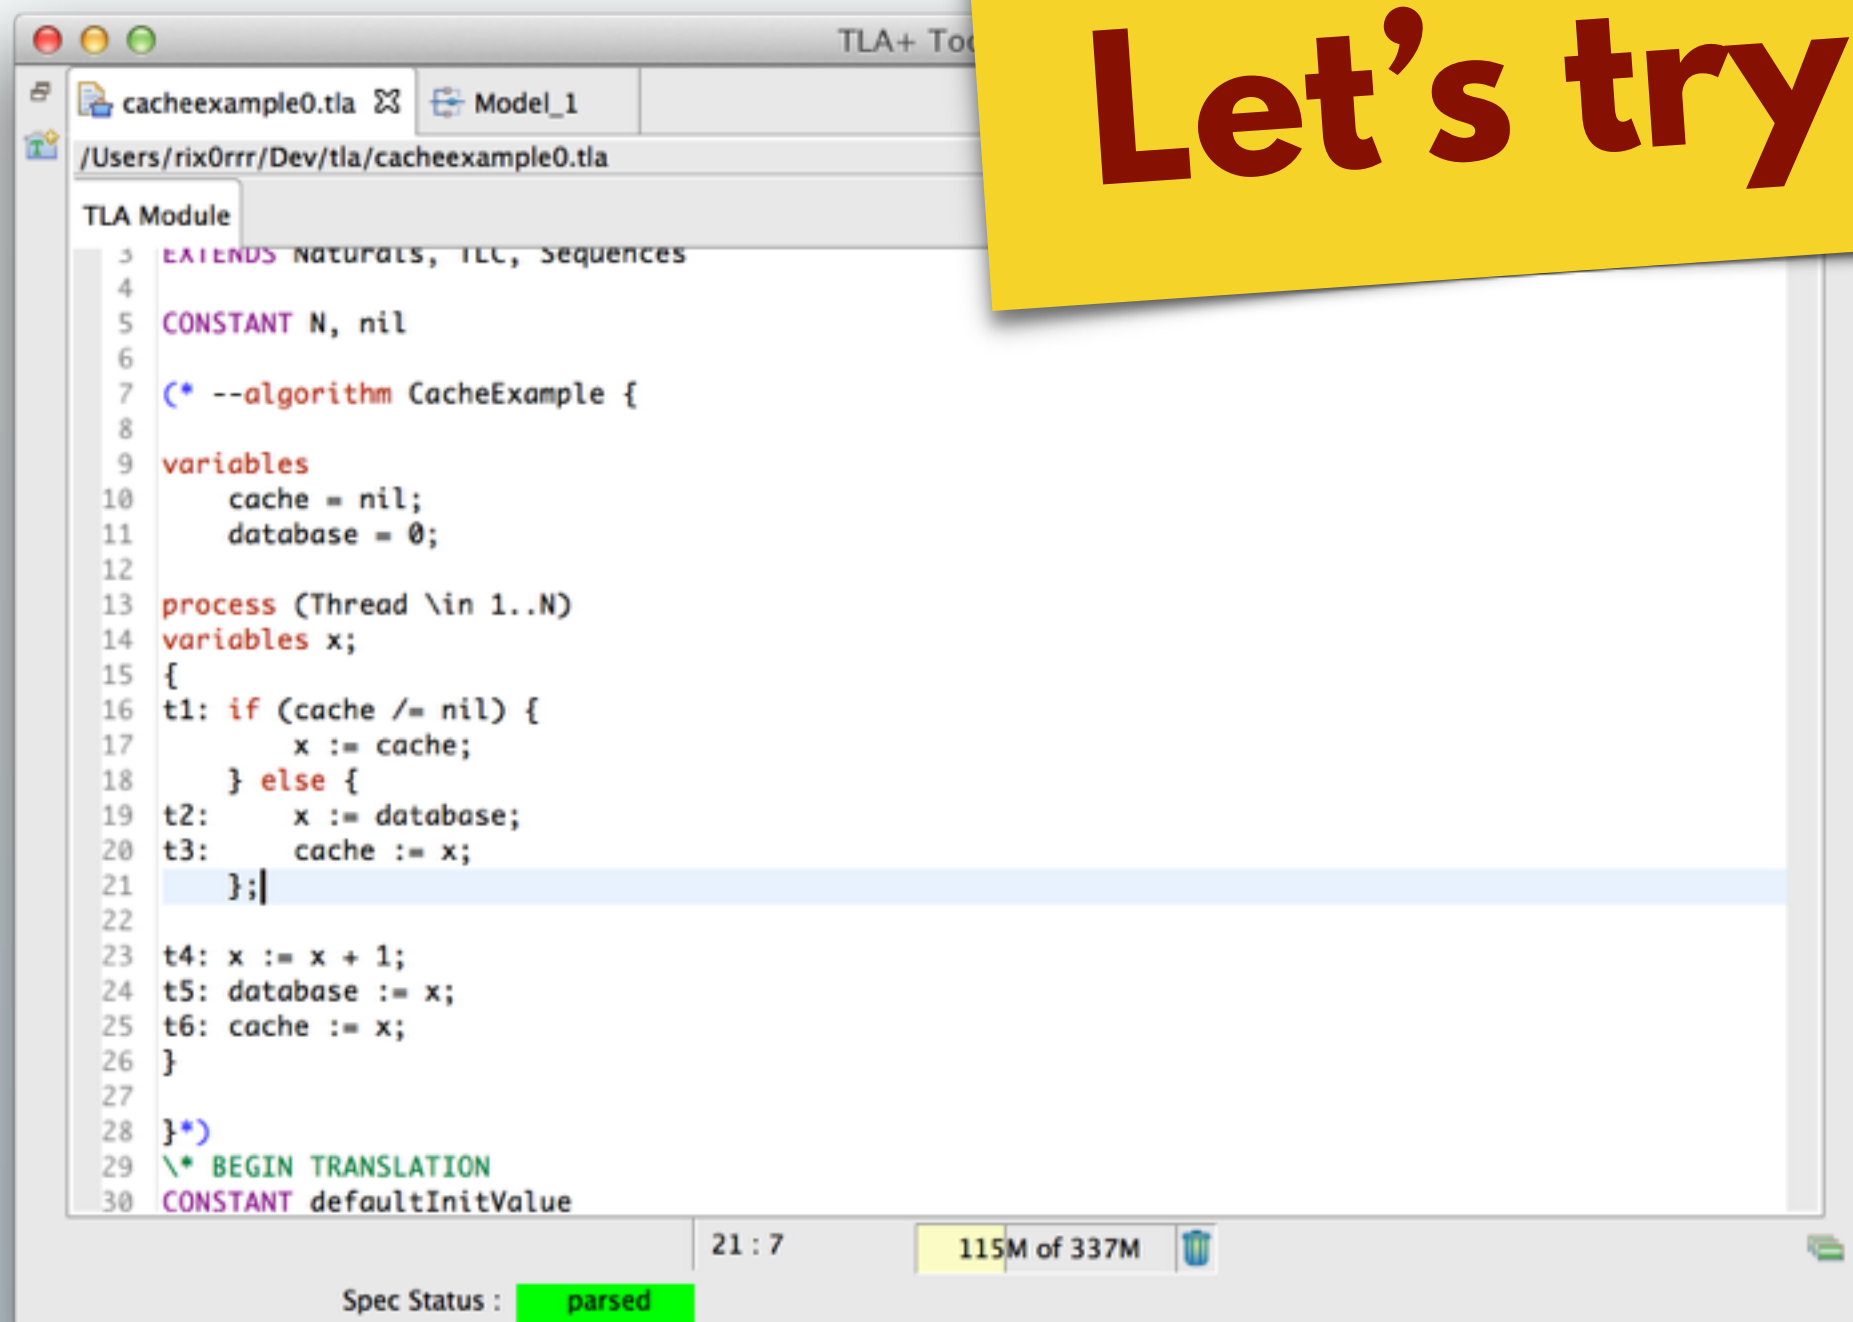
# The Condition?

OK  $\equiv (\forall i \in 1..N : p6[i] = \text{"Done"})$

$\Rightarrow \Rightarrow$

database  $\models N$

Let's try it



The screenshot shows the TLA+ Toolbox editor with a file named 'cacheexample0.tla' open. The editor displays the following TLA code:

```
3 EXTENDS Naturals, TLC, Sequences
4
5 CONSTANT N, nil
6
7 (* --algorithm CacheExample {
8
9   variables
10     cache = nil;
11     database = 0;
12
13   process (Thread \in 1..N)
14   variables x;
15   {
16   t1: if (cache /= nil) {
17       x := cache;
18     } else {
19   t2:   x := database;
20   t3:   cache := x;
21     };|
22
23   t4: x := x + 1;
24   t5: database := x;
25   t6: cache := x;
26   }
27
28   }*)
29 \* BEGIN TRANSLATION
30 CONSTANT defaultInitValue
```

The status bar at the bottom indicates 'Spec Status : parsed' and '115M of 337M'.

<http://rix0rrr.github.io/pluscal-debugger/>



# The Problem?

## Process 1

t=1  
t=2  
t=3  
t=4  
t=5  
t=6

Read

Increment

Write

## Process 2

Read

Increment

Write

**The fix?**

# Conditional Write

```
def Inc(id):  
    x = cache.get(id, None)  
    if not x:  
        x = DB_Read(id)  
        cache[id] = x  
  
    y = x + 1  
    if DB_ConditionalWrite(id, x, y):  
        cache[id] = y  
    else:  
        Inc(id)
```

Exercise: implement in PlusCal!



```

def Inc(id):
    x = cache.get(id, None)
    if not x:
        x = DB_Read(id)
        cache[id] = x

    y = x + 1
    if DB_ConditionalWrite(id, x, y):
        cache[id] = y
    else:
        Inc(id)

```

variables

```

cache = nil;
database = 0;

```

process (Thread \in 1..N)

variables x; y; version;

```

{
t0: if (cache /= nil) {
        x := cache;
    } else {
t1:      x := database;
t2:      cache := x;
    };

```

```

t3: y := x + 1;
t4: if (database = x) {
        database := y;
    } else { goto t0; }; (* Retry *)
t5: cache := y;
}

```





# **Processes instead of threads?**

How would we model that?

variables

database = 0;

cache = nil;

**fair process** (Thread \in 1..N)

variables x; y; version; cache = nil;

{

t0: if (cache /= nil) {

    x := cache;

} else {

t1:     x := database;

t2:     cache := x;

};

t3: y := x + 1;


t4: if (database = x) {

    database := y;

} else { goto t0; }; (\* Retry \*)

t5: cache := y;

}



# Termination

$$Termination \equiv \Diamond (\forall i \in N : pc[i] = \text{“Done”})$$

“at some point in time...”

“all process will be Done”



# General Thoughts

# Writing the spec

- Where the art is
- Minimize state to keep checking feasible. Skip unnecessary detail.
- For a general function that can write to a number of registers, only model *one*
- Don't model actual data *values*, model the *instance* of the data
- Etc.

# Writing the condition

- Write *invariants*
- Write *assertions* inside the code

```
(* Verify that IF the cache isn't hopelessly out of date,  
   THEN the cache is consistent with the database *)  
macro VERIFY_CacheConsistent(t_version) {  
    assert cache_tx_version <= t_version =>  
        \A i \in objs: cache_obj_version[i] \in { 0, db_obj_version[i] };  
}
```

- Maybe *check for termination*

# Common objection



"What if I my code doesn't  
match my model?"



**Thinko-free algorithms**  
**for**  
**very little effort!**