

# Two-Phase Commit

The TLA+ Hyperbook

---

Wang Zhifu 151220117

April 13, 2018

# Table of contents

1. Preliminaries
2. The TLA+ Spec
3. Checking the Spec

# Preliminaries

---

# Records

The definition

$$r \triangleq [prof \mapsto \text{"Fred"}, num \mapsto 42]$$

defines  $r$  to be a record with two fields  $prof$  and  $num$ .

The values of its two fields are

$$r.prof = \text{"Fred"} \text{ and } r.num = 42$$

A record corresponds roughly to a `STRUCT` in C, except that changing the orders of the fields makes no difference.

$$[prof \mapsto \text{"Fred"}, num \mapsto 42] = [num \mapsto 42, prof \mapsto \text{"Fred"}]$$

# Records

The definition

$$r \triangleq [prof \mapsto \text{"Fred"}, num \mapsto 42]$$

defines  $r$  to be a record with two fields  $prof$  and  $num$ .

The values of its two fields are

$$r.prof = \text{"Fred"} \text{ and } r.num = 42$$

A record corresponds roughly to a `STRUCT` in C, except that changing the orders of the fields makes no difference.

$$[prof \mapsto \text{"Fred"}, num \mapsto 42] = [num \mapsto 42, prof \mapsto \text{"Fred"}]$$

# Records

The definition

$$r \triangleq [prof \mapsto \text{"Fred"}, num \mapsto 42]$$

defines  $r$  to be a record with two fields  $prof$  and  $num$ .

The values of its two fields are

$$r.prof = \text{"Fred"} \text{ and } r.num = 42$$

A record corresponds roughly to a `STRUCT` in C, except that changing the orders of the fields makes no difference.

$$[prof \mapsto \text{"Fred"}, num \mapsto 42] = [num \mapsto 42, prof \mapsto \text{"Fred"}]$$

$[prof : \{ "Fred", "Ted", "Ned" \}, num : 0..99]$  is the set of all records.

$[prof \mapsto \dots, num \mapsto \dots]$  with  $prof$  field in  $\{ "Fred", "Ted", "Ned" \}$  and  $num$  field in  $0..99$ .

So  $[prof \mapsto "Ned", num \mapsto 24]$  is in this set.

$[prof : \{ "Fred", "Ted", "Ned" \}, num : 0..99]$  is the set of all records.

$[prof \mapsto \dots, num \mapsto \dots]$  with  $prof$  field in  $\{ "Fred", "Ted", "Ned" \}$  and  $num$  field in  $0..99$ .

So  $[prof \mapsto "Ned", num \mapsto 24]$  is in this set.



$[prof \mapsto "Fred", num \mapsto 42]$  is a function  $f$  with domain  $"prof", "num"$   
s.t.  $f["prof"] = "Fred"$  and  $f["num"] = 42$ .

$f.prof$  is an abbreviation for  $f["prof"]$

$[f \text{ EXCEPT } !["prof"] = "Red"]$

can be abbreviated as

$[f \text{ EXCEPT } !.prof = "Red"]$

$[prof \mapsto \text{"Fred"}, num \mapsto 42]$  is a function  $f$  with domain  $\text{"prof"}, \text{"num"}$   
s.t.  $f[\text{"prof"}] = \text{"Fred"}$  and  $f[\text{"num"}] = 42$ .

$f.\text{prof}$  is an abbreviation for  $f[\text{"prof"}]$

$[f \text{ EXCEPT } ![\text{"prof"}] = \text{"Red"}]$

can be abbreviated as

$[f \text{ EXCEPT } !.\text{prof} = \text{"Red"}]$

$[prof \mapsto \text{"Fred"}, num \mapsto 42]$  is a function  $f$  with domain  $\text{"prof"}, \text{"num"}$   
s.t.  $f[\text{"prof"}] = \text{"Fred"}$  and  $f[\text{"num"}] = 42$ .

$f.\text{prof}$  is an abbreviation for  $f[\text{"prof"}]$

$$[f \text{ EXCEPT } ![\text{"prof"}] = \text{"Red"}]$$

can be abbreviated as

$$[f \text{ EXCEPT } !.\text{prof} = \text{"Red"}]$$

- What Transaction Commit Describes
- TwoPhase Commit Adds the Minister
- A Really Modern Wedding
- A Simplification

# What Transaction Commit Describes

Henry

Anne

# What Transaction Commit Describes



unsure



unsure

# What Transaction Commit Describes



working



working

# What Transaction Commit Describes



prepared



working



# What Transaction Commit Describes



prepared



prepared

# What Transaction Commit Describes



prepared



committed

# What Transaction Commit Describes




committed



committed

## TwoPhase Commit Adds the Minister

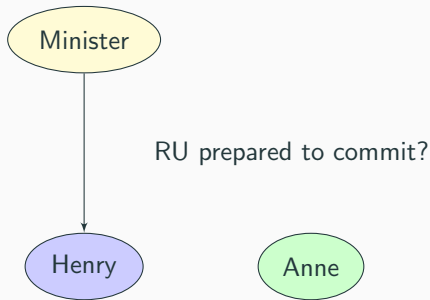


## TwoPhase Commit Adds the Minister


Communication




## TwoPhase Commit Adds the Minister



## TwoPhase Commit Adds the Minister




## TwoPhase Commit Adds the Minister





## TwoPhase Commit Adds the Minister



## TwoPhase Commit Adds the Minister




working




working

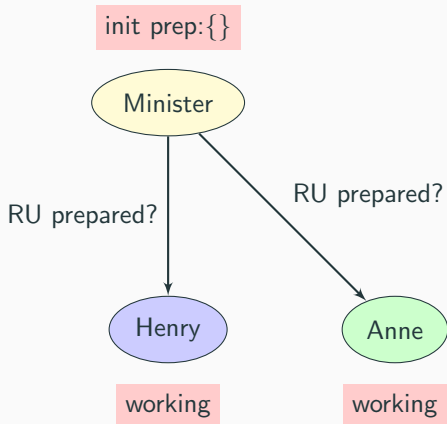
## TwoPhase Commit Adds the Minister



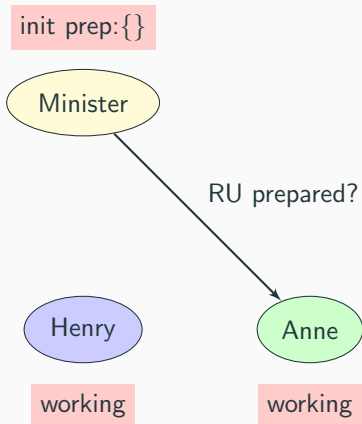
## TwoPhase Commit Adds the Minister



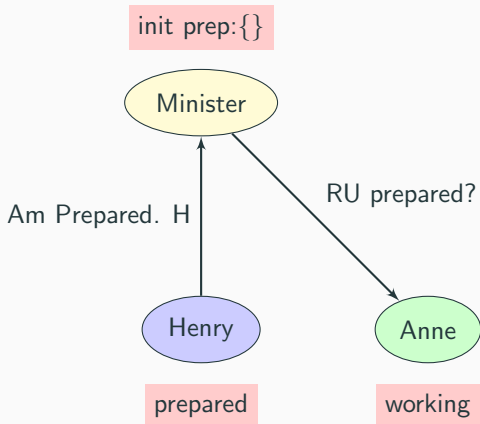
# A Really Modern Wedding



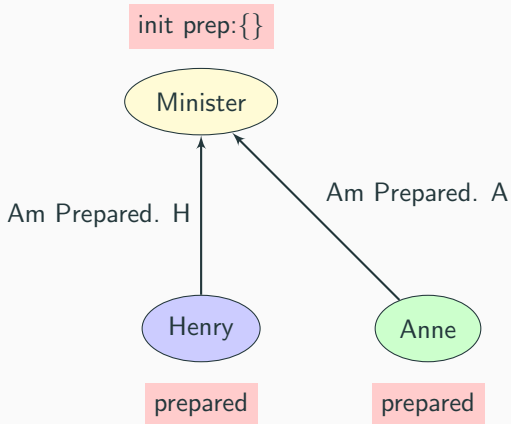
# A Really Modern Wedding



# A Really Modern Wedding

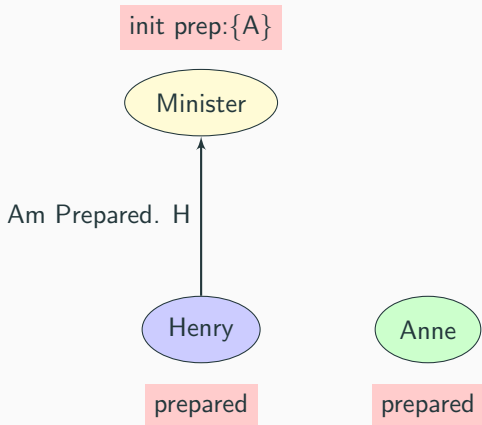


# A Really Modern Wedding





# A Really Modern Wedding



# A Really Modern Wedding

init prep:{A, H}

Minister

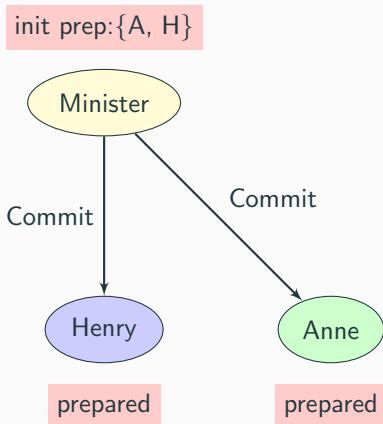
Henry

prepared

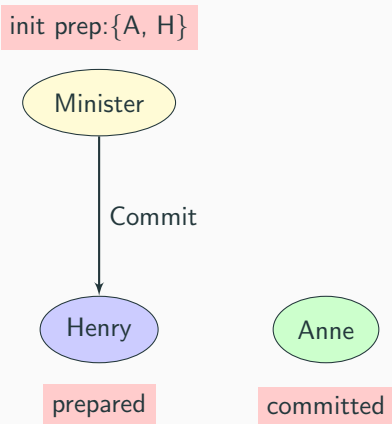
Anne

prepared

# A Really Modern Wedding



# A Really Modern Wedding



# A Really Modern Wedding

init prep: {A, H}

Minister


Henry

committed

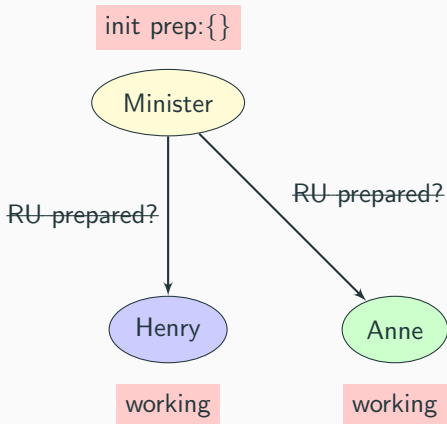
Anne

committed

# A Simplification



# A Simplification



# A Simplification

init prep: {}

Minister

Henry

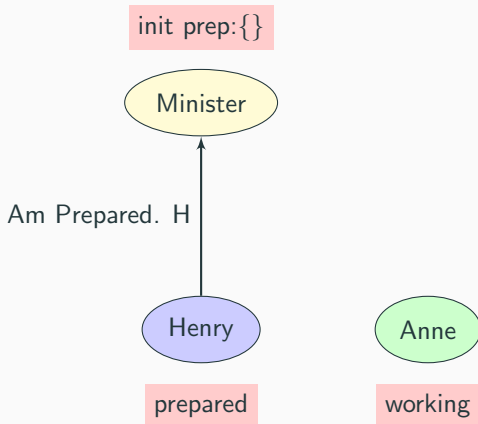
working

Anne

working



# A Simplification



# A Simplification

*RUPrepared?* message not required by *TCommit*.

Simplicity, simplicity, simplicity!

We want the simplest spec that can catch the errors we're looking for, namely, ones that would cause two-phase commit not to satisfy the *TCommit* spec.

# The TLA+ Spec

---

# Declarations of the TLA+ Spec

CONSTANT RM

VARIABLES rmState, tmState, tmPrepared, msgs

*RM*: set of resource managers

*rmState*: state of the resource managers

*tmState*, *tmPrepared*: state of the Transaction Manager

*msgs*: messages that are in transit

# Declarations of the TLA+ Spec

CONSTANT RM

VARIABLES rmState, tmState, tmPrepared, msgs

*RM*: set of resource managers

*rmState*: state of the resource managers

*tmState*, *tmPrepared*: state of the Transaction Manager

*msgs*: messages that are in transit

# Declarations of the TLA+ Spec

CONSTANT RM

VARIABLES rmState, tmState, tmPrepared, msgs

*RM*: set of resource managers

*rmState*: state of the resource managers

*tmState*, *tmPrepared*: state of the Transaction Manager

*msgs*: messages that are in transit

## Sending Messages in the TLA+ Spec

$TPTypeOK \triangleq$

$\wedge rmState \in [RM \rightarrow \{“working”, “prepared”, “committed”, “aborted”\}]$

$\wedge tmState \in \{“init”, “done”\}$

$\wedge tmPrepared \subseteq RM$

$\wedge msgs \subseteq Messages$

## Sending Messages in the TLA+ Spec

It should specify only what's required of message passing.

Let *msgs* be the set of all messages ever sent.

A single message can be received by multiple processes.

A process can receive the same message multiple times.

Two-phase commit still works.



## Sending Messages in the TLA+ Spec

It should specify only what's required of message passing.

Let *msgs* be the set of all messages ever sent.

A single message can be received by multiple processes.

A process can receive the same message multiple times.

Two-phase commit still works.

## Sending Messages in the TLA+ Spec

It should specify only what's required of message passing.

Let *msgs* be the set of all messages ever sent.

A single message can be received by multiple processes.

A process can receive the same message multiple times.

Two-phase commit still works.

## Sending Messages in the TLA+ Spec

It should specify only what's required of message passing.

Let *msgs* be the set of all messages ever sent.

A single message can be received by multiple processes.

A process can receive the same message multiple times.

Two-phase commit still works.

## Sending Messages in the TLA+ Spec

It should specify only what's required of message passing.

Let *msgs* be the set of all messages ever sent.

A single message can be received by multiple processes.

A process can receive the same message multiple times.

Two-phase commit still works.

## Sending Messages in the TLA+ Spec

$TPTypeOK \triangleq$

$\wedge rmState \in [RM \rightarrow \{“working”, “prepared”, “committed”, “aborted”\}]$

$\wedge tmState \in \{“init”, “done”\}$

$\wedge tmPrepared \subseteq RM$

$\wedge msgs \subseteq Messages$

## Sending Messages in the TLA+ Spec

$Messages \triangleq [type : "Prepared", rm : RM] \cup [type : "Commit", "Abort"]$

$[type \mapsto "Prepared", rm \mapsto r]$

represents a Prepared message sent by  $r$  to the TM.

$[type \mapsto "Commit"], [type \mapsto "Abort"]$

Each record represents a message sent by the TM to all RMs.

## Sending Messages in the TLA+ Spec

$TPInit \triangleq$

$\wedge rmState = [r \in RM \mapsto \text{"working"}]$

$\wedge tmState = \text{"init"}$

$\wedge tmPrepared = \{\}$

$\wedge msgs = \{\}$

## Sending Messages in the TLA+ Spec

$$\begin{aligned} \text{TMRCvPrepared}(r) &\triangleq \\ &\wedge \text{tmState} = \text{"init"} \\ &\wedge [\text{type} \mapsto \text{"Prepared"}, \text{rm} \mapsto r] \in \text{msgs} \\ &\wedge \text{tmPrepared}' = \text{tmPrepared} \cup \{r\} \\ &\wedge \text{UNCHANGED } \langle \text{rmState}, \text{tmState}, \text{msgs} \rangle \end{aligned}$$

Enabling conditions: conditions on the first state of a step

All subsequent  $\text{TMRCvPrepared}(r)$  steps leave all the variables unchanged.



# THE REST OF THE SPEC

*TMCommit*  $\triangleq$

It allows steps where the TM sends Commit messages to the RMs and sets `tmState` to “done”.

It is enabled if `tmState` equals “init” and `tmPrepared` equals RM.

```
TMCommit ==  
  /\ tmState = "init"  
  /\ tmPrepared = RM  
  /\ tmState' = "done"  
  /\ msgs' = msgs \cup {[type |-> "Commit"]}  
  /\ UNCHANGED <<rmState, tmPrepared>>
```

# THE REST OF THE SPEC

*TMAbort*  $\triangleq$

The TM sends Abort messages to the RMs and sets tmState to “done”.

It is enabled if tmState equals “init”.

```
TMAbort ==  
  /\ tmState = "init"  
  /\ tmState' = "done"  
  /\ msgs' = msgs \cup {[type |-> "Abort"]}  
  /\ UNCHANGED <<rmState, tmPrepared>>
```

# THE REST OF THE SPEC

$RMPPrepare(r) \triangleq$

RM  $r$  sets its state to “*prepared*” and sends a Prepared message to the TM.

It's enabled if  $rmState[r]$  equals “*working*”.

$RMPPrepare(r) ==$

$\wedge rmState[r] = \text{"working"}$

$\wedge rmState' = [rmState \text{ EXCEPT } ![r] = \text{"prepared"}]$

$\wedge msgs' = msgs \setminus \text{cup } \{[type \rightarrow \text{"Prepared"}, rm \rightarrow r]\}$

$\wedge \text{UNCHANGED } \langle\langle tmState, tmPrepared \rangle\rangle$

*RMChooseToAbort*(*r*)  $\triangleq$

When in its “working” state, RM *r* can go to the “aborted” state.

*RMChooseToAbort*(*r*) ==

/\ *rmState*[*r*] = "working"

/\ *rmState*' = [*rmState* EXCEPT ![*r*] = "aborted"]

/\ UNCHANGED <<*tmState*, *tmPrepared*, *msgs*>>

$RM\text{RcvCommitMsg}(r) \triangleq$

$RM\text{RcvAbortMsg}(r) \triangleq$

RM  $r$  receives a “commit” or “abort” message and sets its state accordingly.

# THE REST OF THE SPEC

```
TPNext ==  
  \/ TCommit \/ TAbort  
  \/ \E r \in RM :  
    TMRcvPrepared(r) \/ RMPprepare(r) \/ RMChooseToAbort(r)  
    \/ RMRcvCommitMsg(r) \/ RMRcvAbortMsg(r)
```

## Checking the Spec

---

- Create a New Model
- Check Your Definitions



Symmetry Sets All RMs are identical/interchangeable.

Suppose  $RM = \{ "r1", "r2", "r3" \}$ .

"r1"  $\leftrightarrow$  "r3" in one possible state yields a possible state.

# Correctness of Two-Phase Commit

We've checked that TypeOK is an invariant of the spec.

We should check that formula *TConsistent* of *TCommit*, which asserts that one RM can't commit and another abort, is also an invariant.

The statement

INSTANCE *TCommit*

imports the definitions from *TCommit* into module *TwoPhase*.

Add the invariant *TConsistent* to your model and have TLC check it.

Two-phase commit doesn't just maintain the invariance of *TConsistent*; it implements the specification of transaction commit.

# Correctness of Two-Phase Commit

We've checked that TypeOK is an invariant of the spec.

We should check that formula *TConsistent* of *TCommit*, which asserts that one RM can't commit and another abort, is also an invariant.

The statement

INSTANCE *TCommit*

imports the definitions from *TCommit* into module *TwoPhase*.

Add the invariant *TConsistent* to your model and have TLC check it.

Two-phase commit doesn't just maintain the invariance of *TConsistent*; it implements the specification of transaction commit.

# Correctness of Two-Phase Commit

We've checked that TypeOK is an invariant of the spec.

We should check that formula *TConsistent* of *TCommit*, which asserts that one RM can't commit and another abort, is also an invariant.

The statement

INSTANCE *TCommit*

imports the definitions from *TCommit* into module *TwoPhase*.

Add the invariant *TConsistent* to your model and have TLC check it.

Two-phase commit doesn't just maintain the invariance of *TConsistent*; it implements the specification of transaction commit.

# Correctness of Two-Phase Commit

We've checked that TypeOK is an invariant of the spec.

We should check that formula *TConsistent* of *TCommit*, which asserts that one RM can't commit and another abort, is also an invariant.

The statement

INSTANCE *TCommit*

imports the definitions from *TCommit* into module *TwoPhase*.

Add the invariant *TConsistent* to your model and have TLC check it.

Two-phase commit doesn't just maintain the invariance of *TConsistent*; it implements the specification of transaction commit.

# Correctness of Two-Phase Commit

We've checked that TypeOK is an invariant of the spec.

We should check that formula *TConsistent* of *TCommit*, which asserts that one RM can't commit and another abort, is also an invariant.

The statement

INSTANCE *TCommit*

imports the definitions from *TCommit* into module *TwoPhase*.

Add the invariant *TConsistent* to your model and have TLC check it.

Two-phase commit doesn't just maintain the invariance of *TConsistent*; it implements the specification of transaction commit.

**Thanks for your listening!**